

Drawing

Motivation

- ▶ We want to render something to the screen
- ▶ This requires us to understand how the GPU handles rendering

Basic Idea

- ▶ We simulate motion by drawing scene many times a second
- ▶ Each time, scene may be moved a bit
- ▶ Gives illusion of smooth continuous motion

Rendering

- ▶ Our first render: We'll just draw a single point

Buffers

- ▶ **Concept: Buffer**
 - ▶ GPU is physically separate from CPU
 - ▶ Some GPU's have separate RAM from main memory
 - ▶ So we need to push data from CPU memory to GPU memory
 - ▶ Buffer = chunk of memory accessible to GPU

Buffers

- ▶ Basic pattern for creating buffer:
 - ▶ `glGenBuffer`
 - ▶ `glBindBuffer`
 - ▶ `glBufferData \`
- ▶ This gen/bind/set pattern is used repeatedly in GL

glGenBuffer

- ▶ GL uses integers to refer to *resources*
 - ▶ Resource = Buffer, texture, ...
- ▶ We need to “generate” a unique buffer ID when we want to create a new buffer
- ▶ **Complication:** GL was designed as a C-based API
 - ▶ So its use looks a little un-Pythonic

glGenBuffer

- ▶ We first need to use a special Python interface: The *array*
 - ▶ `import array`
- ▶ `tmp = array.array("I", [0])`
 - ▶ Create an array that holds integers
 - ▶ Initialize it with a single value equal to zero

glGenBuffer

- ▶ Now we can call the GL function:

```
glGenBuffers(1,tmp)  
buffID = tmp[0]
```

- ▶ The “1” tells how many buffers we want to create

glBindBuffer

- ▶ Our program will have *many* buffers (eventually)
- ▶ How does GL know which one we want to work with?
- ▶ Concept: Binding point
 - ▶ A binding point can be associated with one (and only one) buffer
 - ▶ We *bind* a buffer to a binding point when we want to work with it
 - ▶ We *unbind* the buffer when we're done with it

glBindBuffer

- ▶ Code: `glBindBuffer(GL_ARRAY_BUFFER, buffID)`
 - ▶ `GL_ARRAY_BUFFER` is the binding point we're using
 - ▶ This binding point is for vertex data
 - ▶ Second parameter is the buffer we want to work with

glBufferData

- ▶ Now we can ship data from CPU to GPU
- ▶ We'll start off with a simple 2D point
- ▶ Again, we must use the array type since we need to push a blob of raw memory around

```
tmp = array.array("f", [0,0] ).tobytes()  
glBufferData( GL_ARRAY_BUFFER, len(tmp), tmp, GL_STATIC_DRAW)
```

glBufferData

- ▶ First parameter = which binding point we're using
- ▶ Second = size of data, in bytes
- ▶ Third = the data itself
- ▶ Fourth = usage hint
 - ▶ GL_STATIC_DRAW = We specify the data once, but we'll draw with it many times

Done!

- ▶ We're done with this buffer, so we should unbind it:
`glBindBuffer(GL_ARRAY_BUFFER, 0)`
 - ▶ Zero = special value
 - ▶ Tells GL we want nothing attached to the binding point

Note

- ▶ Buffers are so commonly used that we should factor it all out into a separate, reusable class
- ▶ (Do in class)

```
#Buffer.py
class Buffer:
    def __init__(self, dataAsFloatArray ):
        ...
    def bind(self, bindingPoint):
        ...
```

Using It

- ▶ Suppose we create the buffer like so:
`myBuffer = Buffer(array.array("f",[0,0]))`
- ▶ Now we're ready to use it... Right?
 - ▶ Nope.

Using It

- ▶ A buffer doesn't provide enough information for the GPU to use it
- ▶ Remember, it's just a blob of bytes
- ▶ How does the GPU know how to interpret this blob?
 - ▶ It doesn't!

Using It

- ▶ We know it's a list of 2D float values
- ▶ But the GPU does not
- ▶ So we must tell the GPU how information is organized in the buffer

VAO

- ▶ Concept: Vertex Array Object (VAO)
- ▶ Allows us to tell GPU how one or more buffers are to be used
- ▶ Follows same basic pattern as buffers:
 - ▶ `glGenVertexArrays`
 - ▶ `glBindVertexArray`
 - ▶ Set information
 - ▶ Unbind

glGenVertexArrays

- ▶ Same idea as with buffers: Generate a VAO

```
tmp = array.array("I",[0])  
glGenVertexArrays(1,tmp)  
vao = tmp[0]
```

glBindVertexArray

- ▶ Same idea as with buffers: Bind the VAO so we can work with it:
`glBindVertexArray(vao)`
 - ▶ Tells GL we want to work with this particular vao
 - ▶ There's no binding point specified here – just the vao we want to work with

Data

- ▶ Now we will specify data
- ▶ This is where we tell GL about the buffer's layout and that we want to use the buffer

```
myBuffer.bind(GL_ARRAY_BUFFER)
glEnableVertexAttribArray(0)
glVertexAttribPointer( 0, 2, GL_FLOAT, False, 2*4, 0 )
```

Explanation

- ▶ `myBuffer.bind(GL_ARRAY_BUFFER)`
- ▶ We bind the buffer we want to use
- ▶ Easy!

Explanation

- ▶ `glEnableVertexAttribArray(0)`
- ▶ Then we enable the attribute array
- ▶ *Vertex puller* can grab data from up to ≈ 16 buffers at a time
 - ▶ Think of puller as having 16 “pipes”
 - ▶ Data comes down a pipe and gets used for rendering
- ▶ This function call tells GL pipe zero should be enabled
 - ▶ GL records this fact in the vao

Explanation

- ▶ `glVertexAttribPointer(0, 2, GL_FLOAT, False, 2*4, 0)`
- ▶ Finally, we tell GL where to find the data
 - ▶ First parameter = Which pipe we're talking about
 - ▶ Second = how many items per vertex
 - ▶ Third = what type each item is
 - ▶ Fourth = auto-normalize (this is nearly always false)
 - ▶ Fifth = size of each data item, in bytes
 - ▶ Sixth = where item starts in buffer (byte offset)
- ▶ All this information is recorded in the vao
- ▶ In the background, this also takes currently bound buffer and notes that in vao too

Unbind

- ▶ Finally, we unbind: `glBindVertexArray(0)`
 - ▶ Since we're done with this buffer

Draw

- ▶ OK, *now* we can draw something, right?

Draw

- ▶ OK, *now* we can draw something, right?
- ▶ Nope.

Shader

- ▶ We need to tell GPU what to do with the data
- ▶ So far, we've just told the GPU
 - ▶ What our data is (the raw bytes)
 - ▶ How it's organized

GPU


- ▶ Structure of GPU: Highly parallel computer
 - ▶ Designed to process lots of data at the same time
 - ▶ Modern GPU's: Easily in the *teraflop* range
- ▶ We must instruct GPU how to do its work
- ▶ This requires that we write a small program (“shader”) to direct all that parallel processing

GLSL

- ▶ Shader is written in a C-like language called GLSL
 - ▶ GL Shading Language

Pipeline

- ▶ There are actually two shaders to write
- ▶ GPU pipeline:



Explanation

- ▶ Fixed functionality
 - ▶ Hardwired into silicon
 - ▶ Fast
 - ▶ Limited configurability

Fixed Functionality

- ▶ Vertex puller: Grabs data from one or more buffers
 - ▶ Dispatches to vertex shader
- ▶ Rasterizer
 - ▶ Takes output of VS
 - ▶ Determines which pixels are covered

Shaders

- ▶ This is where most of our attention is focused
- ▶ Two shaders of interest: Vertex shader, Fragment shader

Vertex Shader

- ▶ Runs once per vertex
- ▶ Takes input from the buffers we specified in VAO
- ▶ GL expects it will output *screen space location* of that vertex
- ▶ If we're drawing points, GL also expects it to say how many pixels are covered by that point

Fragment Shader

- ▶ Runs once per pixel that is covered by primitive we're drawing
- ▶ Takes input from the buffers we specified in VAO
- ▶ GL expects it will output *screen space location* of that vertex

Vertex Shader

```
#version 430
layout(location=0) in vec2 position;
void main(){
    gl_Position = vec4( position.xy, -1, 1 );
    gl_PointSize = 1;
}
```

- Explain this in-class...

Fragment Shader

```
#version 430
out vec4 color;
void main(){
    color = vec4(1,1,1,1);
}
```

- Explain this in class

GLSL

- ▶ Quick syntax overview:
 - ▶ Every statement must end with a semicolon
 - ▶ Indentation is not significant
 - ▶ Blocks are denoted with { }
 - ▶ All variables (except builtins) must be *declared* before use
 - ▶ Declaration: type, variable name, semicolon

Shaders

- ▶ Shaders are just ordinary text files
- ▶ We can write with any editor we'd like
 - ▶ But must save as plain text!
- ▶ Customary to save with .vs or .fs suffix, but you don't have to
- ▶ We must write code to load data, push to GPU

Pattern

- ▶ We create a *program object*
 - ▶ Acts like container for our shaders
- ▶ Then: For each shader we:
 - ▶ Create a shader object
 - ▶ Specify shader code
 - ▶ Compile shader code
 - ▶ Attach shader to program object
- ▶ Finally, we *link* the shaders together
- ▶ And now we can *use* the program

Code

- ▶ We encapsulate shader operations into their own class: [Program.py](#)

Draw

- ▶ Can we finally draw something?

Draw

- ▶ Can we finally draw something?
- ▶ Yes!

Drawing

- ▶ To draw:
 - ▶ Make sure our program is active
 - ▶ Clear the screen
 - ▶ Bind the VAO
 - ▶ Call `glDrawArrays()`

Draw

- ▶ `glDrawArrays(GL_POINTS,0,1)`
 - ▶ First = Type of primitive to draw
 - ▶ Points, lines, or triangles
 - ▶ Second = Where in buffer to start (0=at beginning)
 - ▶ Third = How many vertices to draw

Result

Wow!




Explanation

- ▶ GL screen space: $-1 \dots 1$ in x , y , and z
- ▶ z doesn't come into play (yet)
- ▶ We put our point at $(0,0)$, so it's in the center of the screen.

Assignment

- ▶ Create and draw a *star field*
 - ▶ Bonus [+33%]: Make the stars different sizes
 - ▶ This will require a bit of independent research on your part
 - ▶ Hint: `glEnable(GL_PROGRAM_POINT_SIZE)`
 - ▶ Possibilities:
 - ▶ Create a second input for each vertex giving size.
 - ▶ Each VS invocation has access to global `gl_VertexID` (0,1,2,...). Use it.



Sources

- ▶ Khronos Group. OpenGL 4 quick reference card.
<http://www.khronos.org>

Created using L^AT_EX.

Main font: Gentium Book Basic, by Victor Gaultney. See <http://software.sil.org/gentium/>

Monospace font: Source Code Pro, by Paul D. Hunt. See <https://fonts.google.com/specimen/Source+Code+Pro> and <http://sourceforge.net/adobe>

Icons by Ulisse Perusin, Steven Garrity, Lapo Calamandrei, Ryan Collier, Rodney Dawes, Andreas Nilsson, Tuomas Kuosmanen, Garrett LeSage, and Jakub Steiner. See <http://tango-project.org>