


More Textures

Motivation

- ▶ We've seen basic texturing
- ▶ Now we'll examine more advanced texturing strategies

Recall

- ▶ Suppose we have a square with texture coordinates that go from 0...1
- ▶ And suppose we have a checkerboard texture
- ▶ We might get something like this:




Question

- ▶ All our texture coordinates were in range 0...1
- ▶ What if they go outside that range?
- ▶ Alter our code: Make texture coordinates go from 0...2
`tbuff = Buffer(array.array("f", [0,0, 0,2, 2,2, 0,2]))`

Result

- ▶ The texture repeats itself:




Repeat

- ▶ Whether we get repeat is governed by the sampler object
- ▶ Change the sampler constructor:
`glSamplerParameteri(self.samp, GL_TEXTURE_WRAP_S,
GL_CLAMP_TO_EDGE)`
`glSamplerParameteri(self.samp, GL_TEXTURE_WRAP_T,
GL_CLAMP_TO_EDGE)`


Result

- ▶ Anywhere the coordinates would have been less than 0 or greater than 1: get *clamped*
- ▶ Remember: (0,0) is lower left corner




Now

- ▶ Suppose we change the position values so they aren't -1...1 anymore
- ▶ Maybe try -0.5...0.5
 - ▶ And while we're at it, make texture coords go 0...4
 - ▶ And turn repeat back on: Use GL_REPEAT instead of GL_CLAMP_TO_EDGE in sampler setup
- ▶ $p=0.5$
`vbuff = Buffer(array.array("f", [-p,-p, p,-p, p,p, -p,p]))`
`t=4`
`tbuff = Buffer(array.array("f", [0,0, 0,t, t,t, 0,t]))`




And

- ▶ Now make $p=0.25$



So...

- ▶ Let's re-set the positions to be $-0.5 \dots 0.5$ and make the texture coordinates go $0 \dots 8$





Notice

- ▶ If we make the object smaller, the checks get smaller
- ▶ Or, if we make the texture coordinates larger, the checks get smaller
- ▶ Maybe it's obvious when you think about it...


Now

- Suppose we bump t to 16. And then to 32.




Well...

- ▶ It seems the limit would be to set $t=64$
- ▶ Then we have one black pixel, one white pixel, one black pixel, ...




Question

- ▶ What happens if we set $t > 64$?
- ▶ How about $t=130$?




Hmm...

- ▶ What if $t=128$?



Or...



- ▶ Maybe $t=90$?



- ▶ What is going on?



Explanation

- ▶ Consider single scan line.
- ▶ Its colors alternate from black to white.
- ▶ We can graph intensities (y axis) versus position (x axis)





Explanation

- ▶ Camera is close to object (or texture coordinates' range is small)
→ Dense sampling
- ▶ Image looks OK





Explanation

- ▶ Camera moves further away (or texture coordinates' range gets larger) → Not as dense sampling
- ▶ Reconstruction still OK.





Explanation

- ▶ Camera is further away (or texture coordinates' range gets very large) → Sparse sampling
- ▶ Checks are uneven now



Explanation

- ▶ Camera is even further away (or texture coordinates' range very very large) → Very sparse sampling
- ▶ “Large check” pattern
- ▶ This is *aliasing*
 - ▶ Name comes from signal theory



Human Eyes

- ▶ Two factors help prevent aliasing in human eyes:
 - ▶ Photoreceptors are not in regular grid pattern
 - ▶ Regular patterns make aliasing more prominent
 - ▶ Cornea acts as “low pass filter”
 - ▶ Blurs out high frequency components

Solution

- ▶ We can't really change the grid pattern of monitor pixels
- ▶ But we can simulate the cornea's action
 - ▶ The further something is, the more we blur it

How?

- ▶ How to accomplish blurring?
 - ▶ Too time consuming to do on the fly
 - ▶ How can we quantify amount of blur to perform?

Mip Mapping

- ▶ Mip mapping: Introduced in 1983 by Lance Williams
 - ▶ Mip = Multum in Parvo = Many things in a small place
- ▶ Preprocess the texture to get successively more blurred versions of it
- ▶ Store them all
- ▶ At runtime, select one based on distance from viewer
 - ▶ Further away = More blurred
- ▶ The details are a bit subtle...

Question

- ▶ How to determine distance from viewer?
- ▶ GPU uses a heuristic: For each 2x2 quad of pixels:
 - ▶ Look at texel coordinates accessed for each quad
 - ▶ Difference between them allows us to estimate size of a pixel in texel space
- ▶ The bigger the pixel is in texel space, the further away the object must be

MIP

- ▶ Bigger pixel size → Biases the selection to a more blurred copy of image
- ▶ Also referred to as a *higher mip level*

Terminology


- ▶ Mip level 0 = the original, full detail image
- ▶ Mip level 1 = a little blurred
- ▶ Mip level 2 = more blurred
- ▶ etc.

Blurred Images

- ▶ One problem: Storing many copies of image chews up GPU memory
 - ▶ We will likely have *many* textures and models competing for space
 - ▶ So we don't want to waste RAM needlessly

Solution

- ▶ Suppose we have mip level 0: Our original image. Suppose it has size 128x48
 - ▶ Take four adjacent pixels from mip level 0 and average them together to get one pixel
 - ▶ Do this for all the pixels of level 0
 - ▶ This gives us an image that is half the size of level 0 (64x24)
 - ▶ We'll refer to this as mip level 1



Solution

- ▶ Repeat this process to get mip level 2 (32x12)
- ▶ And again for level 3 (16x6)
- ▶ And once more for level 4 (8x3)
- ▶ One more time for level 5 (4x1)
 - ▶ Notice: $3/2$ gives a noninteger quotient
 - ▶ Power-of-two dimensions work best for mipmapping, but modern GPU's can handle NPOT textures
- ▶ Level 6 will be 2x1
- ▶ And level 7 is 1x1. We're done.

Memory

- ▶ We said memory usage was a concern.
- ▶ How much RAM does a mipmapped image require?
- ▶ Suppose input is 32x32. RAM: $32 \times 32 \times 3 = 3072$ bytes
- ▶ Level 1 = $16 \times 16 \times 3 = 768$
- ▶ Level 2 = $8 \times 8 \times 3 = 192$
- ▶ Level 3 = $4 \times 4 \times 3 = 48$
- ▶ Level 4 = $2 \times 2 \times 3 = 12$
- ▶ Level 5 = $1 \times 1 \times 3 = 3$
- ▶ Total: 4096 bytes
 - ▶ Which is 133% of original 3072 bytes
- ▶ Mip mapping adds overhead of 33%.

Result

- ▶ How do we use mipmapping?
- ▶ GL includes it built-in
- ▶ We must request it from our sampler: Change MAG filter to GL_LINEAR and MIN filter to GL_LINEAR_MIPMAP_LINEAR

And

- ▶ In `ImageTexture2DArray`: After the `glTexImage3D` call:
`glGenerateMipmap(GL_TEXTURE_2D_ARRAY)`
- ▶ This is important: Without this, the mip chain is *incomplete*
- ▶ GL will always return black (0,0,0,0) for texture with incomplete mip chain
- ▶ My GPU gives a debug message:
131204 : Texture state usage warning: The texture object (1) bound to texture image unit 0 does not have a complete set of mipmaps and cannot be used with a sampler needing mipmaps.


Result

- ▶ With $t=90$: The checkers are starting to blur out to grey





Result

- ▶ With $t=130$: Notice it's now a uniform grey color, simulating a blurred checkerboard




Texel Shape

- ▶ Recall how we said we'd select mip level: Look at difference between texture coordinates to estimate pixel size in texture space
- ▶ But: A pixel is usually *not* a square in texture space
 - ▶ What mipmap level to use? s says one thing, t says another
 - ▶ Could average them, prefer maximum one, prefer minimum one, etc.





Anisotropic

- ▶ Another strategy:
 - ▶ Map four corners of pixel to texture space
 - ▶ Take several samples along longest axis
 - ▶ Average them together
- ▶ Reduces blurring, especially in perspective views
- ▶ Slower (but has hardware support)



Anisotropic

- ▶ Left: Linear mipmap filtering
- ▶ Right: Anisotropic filtering



To Use

- ▶ To use anisotropic filtering: Set `GL_TEXTURE_MAX_ANISOTROPY_EXT` sampler parameter
 - ▶ Value is an integer: 1=no anisotropy, 2=use 2 samples, 3=use 3 samples, ...

Assignment

- ▶ Enable mipmaps
- ▶ Add enemies that spawn at random times and move right-to-left across the screen. When an enemy goes off the left side, it should be deleted
- ▶ Your enemy should have a unique texture.
- ▶ For bonus [+25], have two types of enemies (with different textures): One which moves right-to-left and the other drops down from above.

Sources

- ▶ A. Watt & M. Watt. *Advanced Animation and Rendering Techniques*. Addison-Wesley.
- ▶ Sampler Object. OpenGL Wiki.
https://www.khronos.org/opengl/wiki/Sampler_Object

Created using L^AT_EX.

Main font: Gentium Book Basic, by Victor Gaultney. See <http://software.sil.org/gentium/>

Monospace font: Source Code Pro, by Paul D. Hunt. See <https://fonts.google.com/specimen/Source+Code+Pro> and <http://sourceforge.net/adobe>

Icons by Ulisse Perusin, Steven Garrity, Lapo Calamandrei, Ryan Collier, Rodney Dawes, Andreas Nilsson, Tuomas Kuosmanen, Garrett LeSage, and Jakub Steiner. See <http://tango-project.org>